

UNIX is basically a simple operating system, but you have to be a genius to understand the simplicity.

Dennis Ritchie

2

Phase 1 - Level 2: The Queues Manager

Level 2 of Pandos instantiates the key operating system concept that active entities at one layer are just data structures at lower layers. In this case, the active entities at a higher level are processes (i.e. programs in execution) and the data structure(s) that represent them at this level are *process control blocks* (*pcbs*).

```

/* process control block type */
typedef struct pcb_t {
    /* process queue fields */
    struct pcb_t      *p_next,      /* pointer to next entry */
                      *p_prev,      /* pointer to prev entry */

    /* process tree fields */
                      *p_prnt,      /* pointer to parent      */
                      *p_child,     /* pointer to 1st child  */
                      *p_sib;       /* pointer to sibling     */

    /* process status information */
    state_t           p_s;          /* processor state */
    cpu_t             p_time;       /* cpu time used by proc */
    int               *p_semAdd;    /* pointer to sema4 on  */
                                /* which process blocked */

    /* support layer information */
    support_t         *p_supportStruct;
                                /* ptr to support struct */

    /* namespace layer information */
    nsd_t             *namespaces[NS_TYPE_MAX];
                                /* Namespace pointer per type */
} pcb_t;

```

The queue manager will implement four *pcb* related sets of functions:

- The allocation and deallocation of *pcbs*.
- The maintenance of queues of *pcbs*.
- The maintenance of trees of *pcbs*.
- The maintenance of an hash table of *active semaphore descriptors*, each of which supports a queue of *pcbs*: The ASH.

- The maintenance of a list of *namespace descriptors* per type.

2.1 The Allocation and Deallocation of *pcbs*

One may assume that Pandos supports no more than *MAXPROC* concurrent processes; where *MAXPROC* should be set to 20 (in the `const.h`) file.¹ Thus this level needs a “pool” of *MAXPROC pcbs* to allocate from and deallocate to. Assuming that there is a set of *MAXPROC pcbs*, the free or unused ones can be kept on a NULL-terminated single, linearly linked list (using the `p_next` field), called the *pcbFree* List, whose head is pointed to by the variable `pcbFree_h`.

To support the allocation and deallocation of *pcbs* there should be the following three externally visible functions:

- *pcbs* which are no longer in use can be returned to the `pcbFree` list by using the method:

```
void freePcb(pcb_t *p)
```

```
    /* Insert the element pointed to by p onto the pcbFree list.
    */
```

- *pcbs* should be allocated by using:

```
pcb_t *allocPcb()
```

```
    /* Return NULL if the pcbFree list is empty. Otherwise, re-
    move an element from the pcbFree list, provide initial values
    for ALL of the pcbs fields (i.e. NULL and/or 0) and then
    return a pointer to the removed element. pcbs get reused, so
    it is important that no previous value persist in a pcb when
    it gets reallocated. */
```

There is still the question of how one acquires storage for *MAXPROC pcbs* and gets these *MAXPROC pcbs* initially onto the `pcbFree` list. Unfortunately, there is no `malloc()` feature to acquire dynamic (i.e. non-automatic) storage that will persist for the lifetime of the OS and not just the lifetime of the function they are declared in. Instead, the storage for the *MAXPROC pcbs* will be allocated as *static* storage. A static array of *MAXPROC pcbs* will

¹A supplied “starter” version of `const.h` can be found in `usr/include/umps3/umps`

be declared in `initPcb()`. Furthermore, this method will insert each of the MAXPROC *pcbs* onto the `pcbFree` list.

- To initialize the `pcbFree` List:

```
initPcb()
```

```
/* Initialize the pcbFree list to contain all the elements of
the static array of MAXPROC pcbs. This method will be
called only once during data structure initialization. */
```

2.2 Process Queue Maintenance

The methods below do not manipulate a particular queue or set of queues. Instead they are generic queue manipulation methods; one of the parameters is a pointer to the queue upon which the indicated operation is to be performed.

The queues of *pcbs* to be manipulated, which are called *process queues*, are all double, circularly linked lists, via the `p_next` and `p_prev` pointer fields. Instead of a head pointer, each queue will be pointed at by a tail pointer.

To support process queues there should be the following externally visible functions:

```
pcb_t *mkEmptyProcQ()
```

```
/* This method is used to initialize a variable to be tail pointer
to a process queue.
Return a pointer to the tail of an empty process queue; i.e. NULL.
*/
```

```
int emptyProcQ(pcb_t *tp)
```

```
/* Return TRUE if the queue whose tail is pointed to by tp is
empty. Return FALSE otherwise. */
```

```
insertProcQ(pcb_t **tp, pcb_t *p)
```

```
/* Insert the pcb pointed to by p into the process queue whose tail-
pointer is pointed to by tp. Note the double indirection through
tp to allow for the possible updating of the tail pointer as well.
*/
```

```
pcb_t *removeProcQ(pcb_t **tp)
```

```
/* Remove the first (i.e. head) element from the process queue
whose tail-pointer is pointed to by tp. Return NULL if the process
queue was initially empty; otherwise return the pointer to
the removed element. Update the process queue's tail pointer if
necessary. */
```

```
pcb_t *outProcQ(pcb_t **tp, pcb_t *p)
```

```
/* Remove the pcb pointed to by p from the process queue whose
tail-pointer is pointed to by tp. Update the process queue's tail
pointer if necessary. If the desired entry is not in the indicated
queue (an error condition), return NULL; otherwise, return p.
Note that p can point to any element of the process queue. */
```

```
pcb_t *headProcQ(pcb_t *tp)
```

```
/* Return a pointer to the first pcb from the process queue whose
tail is pointed to by tp. Do not remove this pcb from the process
queue. Return NULL if the process queue is empty. */
```

2.3 Process Tree Maintenance

In addition to possibly participating in a process queue, *pcbs* are also organized into trees of *pcbs*, called *process trees*. The `p_prnt`, `p_child`, and `p_sib` pointers are used for this purpose.

The process trees should be implemented as follows. A parent *pcb* contains a pointer (`p_child`) to a NULL-terminated single, linearly linked list of its child *pcbs*. Each child process has a pointer to its parent *pcb* (`p_prnt`) and possibly the next child *pcb* of its parent (`p_sib`). For greater efficiency you may want to make the linked list of child *pcbs* a NULL-terminated double, linearly linked list.

To support process trees there should be the following externally visible functions:

```
int emptyChild(pcb_t *p)
```

```
/* Return TRUE if the pcb pointed to by p has no children.
Return FALSE otherwise. */
```

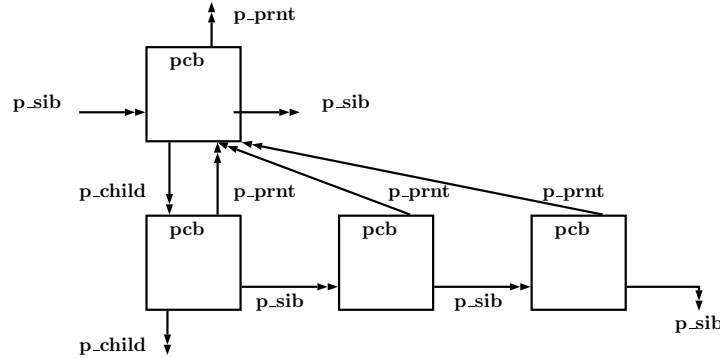


Figure 2.1: Process Tree

```
insertChild(pcb_t *prnt, pcb_t *p)
```

```
/* Make the pcb pointed to by p a child of the pcb pointed to by
prnt. */
```

```
pcb_t *removeChild(pcb_t *p)
```

```
/* Make the first child of the pcb pointed to by p no longer a
child of p. Return NULL if initially there were no children of p.
Otherwise, return a pointer to this removed first child pcb. */
```

```
pcb_t *outChild(pcb_t *p)
```

```
/* Make the pcb pointed to by p no longer the child of its parent.
If the pcb pointed to by p has no parent, return NULL; otherwise,
return p. Note that the element pointed to by p need not be the
first child of its parent. */
```

2.4 The Active Semaphore Hash (ASH)

A *semaphore* is an important operating system concept. While understanding semaphores is not yet needed, this level nevertheless implements an important data structure/abstraction which supports Pandos's implementation of semaphores.

For the purpose of this level it is sufficient to think of a semaphore as an integer. Associated with this integer is:

- An address; semaphores, like all integers, have a physical address in memory.
- A process queue.

A semaphore is *active* if there is at least one *pcb* on the process queue associated with it. (i.e. The process queue is not empty: `emptyProcQ(s_procq)` is FALSE.)

The following implementation is suggested: Maintain an hash table of semaphore descriptors which key is represented by its pointer. The hash is represented by `semd_h`. The hash `semd_h` points to will represent the *Active Semaphore Hash* (ASH).

```
/* semaphore descriptor type */
typedef struct semd_t {
    struct hash_table_entry s_link;    /* ASH reference */
    struct semd_t          *s_freelink; /* next element on the */
                                   /* free semaphore list */
    int                    *s_semAdd;   /* semaphore pointer */
    pcb_t                  *s_procQ;   /* tail pointer to a */
                                   /* process queue */
} semd_t;
```

Maintain a list of semaphore descriptors, the *semdFree* list, to hold the unused semaphore descriptors. This list, whose head is pointed to by the variable `semdFree_h`, is kept, like the *pcbFree* list, as a NULL-terminated single, linearly linked list (using the `s_freelink` field).

The semaphore descriptors themselves should be declared, like the *pcbs*, as a static array of size `MAXPROC` of type `semd_t`.

To support the ASH there should be the following externally visible functions:

```
int insertBlocked(int *semAdd, pcb_t *p)
```

```
/* Insert the pcb pointed to by p at the tail of the process queue
associated with the semaphore whose physical address is semAdd
and set the semaphore address of p to semAdd. If the semaphore
is currently not active (i.e. there is no descriptor for it in the
ASH), allocate a new descriptor from the semdFree list, insert it
in the ASH, initialize all of the fields (i.e. set s_semAdd to semAdd,
```

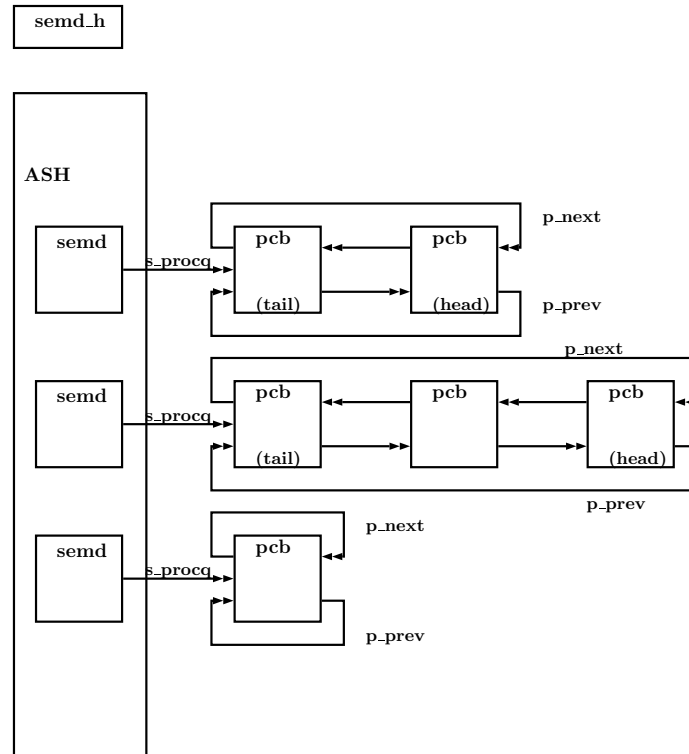


Figure 2.2: Active Semaphore Hash

and `s_procq` to `mkEmptyProcQ()`), and proceed as above. If a new semaphore descriptor needs to be allocated and the `semFree` list is empty, return `TRUE`. In all other cases return `FALSE`. */

```
pcb_t *removeBlocked(int *semAdd)
```

```
/* Search the ASH for a descriptor of this semaphore. If none is
found, return NULL; otherwise, remove the first (i.e. head) pcb
from the process queue of the found semaphore descriptor and
return a pointer to it. If the process queue for this semaphore
becomes empty (emptyProcQ(s_procq) is TRUE), remove the
semaphore descriptor from the ASH and return it to the semFree
list. */
```

```
pcb_t *outBlocked(pcb_t *p)
```



```

/* Remove the pcb pointed to by p from the process queue as-
sociated with p's semaphore (p→ p_semAdd) on the ASH. If pcb
pointed to by p does not appear in the process queue associated
with p's semaphore, which is an error condition, return NULL;
otherwise, return p. */

pcb_t *headBlocked(int *semAdd)
/* Return a pointer to the pcb that is at the head of the process
queue associated with the semaphore semAdd. Return NULL if
semAdd is not found on the ASH or if the process queue associated
with semAdd is empty. */

initASH()
/* Initialize the semFree list to contain all the elements of the
array
static sem_t semTable[MAXPROC]
This method will be only called once during data structure ini-
tialization. */

```

Technical Point: Strive to structure the ASH code so that there is one internal/helper function that traverses the ASH and is used by `insertBlocked`, `removeBlocked`, `outBlocked`, and `headBlocked`.

2.5 Namespace management

A *namespace* is an operating system-level virtualization concept. Every process or thread can have its view on the Pandos's system.

Every namespace is then represented by a type implemented as an integer.

A NULL value in the namespace key represent the *base* namespace. i.e. the namespace in which, by default, every process resides.

You can use the following structure to describe a namespace.

```

/* namespace descriptor type */
typedef struct nsd_t {
    struct nsd_t *n_next;    /* next element on the free list */
    int          n_type;     /* pointer to the semaphore */
} nsd_t;

```

Maintain a list of free namespace descriptors per each type, e.g. *pid_nsFree* indexed by their pointer to hold unused namespace descriptors. These lists, whose are represented by the variables **type_nsFree**, are kept, like the ASH free list, in various global variables.

The namespace descriptors themselves should be declared, like the *pcbs*, as one array of size MAXPROC (per namespace type) of type **nsd_t**. For the sake of phase1, only a single namespace type (PID_NS) should be declared.

To support the Namespaces there should be the following externally visible functions:

```
nsd_t *getNamespace(pcb_t *p, int type)

    /* Return the pointer to the namespace descriptor of type type
    associated with the pcb p. */

nsd_t *allocNamespace(int type)

    /* Allocate a namespace from the type_nsFree list and return
    to the user, this value can be used for the next calls to refer to
    this namespace. */

void freeNamespace(nsd_t *descriptor)

    /* Free a namespace, adding its list pointer (n.link) to the cor-
    rect type_nsFree list. */

int addNamespace( *ns, pcb_t *p)

    /* Insert the namespace ns as the namespace for the correct type
    of p to ns. If the namespace is currently the base (i.e. there are
    no descriptor for the namespace), allocate a new descriptor from
    the type_nsFree list, insert it in the correct list, initialize all of
    the fields, and proceed as above. If a new namespace descriptor
    needs to be allocated and the type_nsFree list is empty, return
    TRUE. In all other cases return FALSE. */

void initNamespaces()

    /* Initialize the type_nsFree list to contain all the elements of the
    arrays
    static ns_t type_nsTable[MAXPROC]
    This method will be only called once during data structure ini-
    tialization. */
```

2.6 Nuts and Bolts

There is no one right way to implement the functionality of this level. The recommended approach is to create three modules (i.e. files): one for the ASH, one for the Namespace management, and one for *pcb* initialization, allocation, deallocation, process queue maintenance, and process tree maintenance.

The second module, `pcb.c`, in addition to the public and `HIDDEN`/private helper functions, will also contain the declaration for the private global variable that points to the head of the `pcbFree` list.

```
HIDDEN pcb_t *pcbFree_h;
```

The ASL module, `asl.c`, in addition to the public and `HIDDEN`/private helper functions, will also contain the declarations for `semd.h` and `semdFree.h`

```
HIDDEN semd_t *semd_h, *semdFree_h;
```

Since the ASL module will make calls to the process queue module to manipulate the process queue associated with each active semaphore, this module should

```
#include "pcb.h"
```

This will insure that the ASH can only use the externally visible functions from `pcb.c` for maintaining its process queues.

Furthermore, the declaration for `pcb_t` would then be placed in the `types.h` file.² This is because many other modules will need to access this definition. The declaration for `semd_t` can be placed in either `asl.c` (because no other module will ever need to access this definition), or `types.h`.

2.7 Testing

There is a provided test file, `p1test.c` that will “exercise” your code. [Appendix ??]

²A supplied “starter” version of `types.h` can be found in `/usr/include/umps3/umps`

As with any non-trivial system, you are strongly encouraged to use the *make* program to maintain your code. A sample *Makefile* has been supplied for you to use. See Chapter ?? in the POPS reference for more compilation details.

Once your (three?) source files have been correctly compiled, linked together (with appropriate linker script, `crtso.o`, and `libumps.o`), and post-processed with `umps3-elf2umps` (all performed by the sample *Makefile*), your code can be tested by launching the μ MPS3 emulator. At a terminal prompt, enter:

```
umps3
```

The test program reports on its progress by writing messages to `TERMINAL0`. These messages are also added to one of two memory buffers; `errbuf` for error messages and `okbuf` for all other messages. At the conclusion of the test program, either successful or unsuccessful, μ MPS3 will display a final message and then enter an infinite loop. The final message will either be **System Halted** for successful termination, or **Kernel Panic** for unsuccessful termination.